

Other Generalized Linear Models

Logistic Binomial Model For count data we have discussed Poisson and Negative-Binomial sampling models. *It is also possible to use a Binomial distribution, but know that the support of the response will not be countably infinite.*

A common example of binomial data would be free throw shooting for basketball players or batting data for baseball players.

```
batting <- read_csv('http://math.montana.edu/ahoegh/teaching/stat491/data/BattingAverage.csv') %>%
  mutate(NotHits = AtBats - Hits)
```

```
## Parsed with column specification:
## cols(
##   Player = col_character(),
##   PriPos = col_character(),
##   Hits = col_double(),
##   AtBats = col_double(),
##   PlayerNumber = col_double(),
##   PriPosNumber = col_double()
## )
```

```
batting %>% sample_n(5)
```

```
## # A tibble: 5 x 7
##   Player      PriPos      Hits AtBats PlayerNumber PriPosNumber NotHits
##   <chr>      <chr>    <dbl> <dbl>      <dbl>      <dbl>      <dbl>
## 1 Martin Maldonado Catcher      62   233         539          2     171
## 2 Jeff Suppan    Pitcher       1    10         842          1       9
## 3 Nathan Eovaldi Pitcher       3    32         252          1     29
## 4 Eric Young     Center Field  55   174         940          8    119
## 5 Reed Johnson   Left Field   78   269         445          7    191
```

The logistic-binomial framework is written as:

$$y_i \sim \text{Binomial}(n_i, p_i), \quad (1)$$

$$\text{logit}(p_i) = X_i \beta \quad (2)$$

```
log_binom <- stan_glm(cbind(Hits, NotHits) ~ PriPos - 1,
  family = binomial(link = "logit"), data = batting, refresh = 0)

print(log_binom, digits = 2)
```

```
## stan_glm
## family:      binomial [logit]
## formula:     cbind(Hits, NotHits) ~ PriPos - 1
## observations: 948
## predictors:  9
## -----
##              Median MAD_SD
## PriPos1st Base    -1.05  0.02
## PriPos2nd Base    -1.07  0.02
## PriPos3rd Base    -1.02  0.02
## PriPosCatcher     -1.11  0.02
## PriPosCenter Field -1.03  0.02
## PriPosLeft Field  -1.05  0.02
## PriPosPitcher     -1.91  0.04
## PriPosRight Field -1.03  0.02
## PriPosShortstop   -1.07  0.02
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

Overdispersion can also occur with binomial data. Recall that the variance of binomial trials is $np(1-p)$. Then define the residuals as $z_i = \frac{y_i - \hat{y}_i}{\text{sd}(\hat{y}_i)}$.

Then the z_i terms should be approximately iid $N(0,1)$. A formal test for $\sum z_i^2$ using a χ^2 distribution can be used to detect overdispersion.

Often hierarchical models will solve some of these issues, otherwise an overdispersion model can be formulated with variance equal to $\omega np(1-p)$. See `brm` or write your own in `stan`.

Probit Model Consider an alternative link function for binary/binomial data.


$$y_i \sim \text{Binomial}(n_i, p_i), \quad (3)$$

$$\Phi^{-1}(p_i) = X_i \beta \quad (4)$$

$$p_i = \Phi(X_i \beta), \quad (5)$$

where $\Phi()$ is the cumulative distribution function for a standard normal random variable.

This model is a latent data model, which are very common and useful in statistics. We assume there is an underlying continuous random variable that is mapped to a standard normal distribution.



where

$$y_i = \begin{cases} 1 & \text{if } z_i > 0 \\ 0 & \text{if } z_i < 0 \end{cases}$$

$$z_i = X_i \beta + \epsilon$$


$$\epsilon \sim N(0, 1)$$

Note $\epsilon \sim N(0, 1)$ is a necessary constraint for this model.

```
probit_binom <- stan_glm(cbind(Hits, NotHits) ~ PriPos - 1,
  family = binomial(link = "probit"), data = batting, refresh = 0)
print(probit_binom)
```

```
## stan_glm
## family:      binomial [probit]
## formula:     cbind(Hits, NotHits) ~ PriPos - 1
## observations: 948
## predictors:  9
## -----
##              Median MAD_SD
## PriPos1st Base   -0.6    0.0
## PriPos2nd Base   -0.7    0.0
## PriPos3rd Base   -0.6    0.0
## PriPosCatcher    -0.7    0.0
## PriPosCenter Field -0.6    0.0
## PriPosLeft Field  -0.6    0.0
## PriPosPitcher    -1.1    0.0
## PriPosRight Field -0.6    0.0
## PriPosShortstop  -0.7    0.0
##
## -----
## * For help interpreting the printed output see ?print.stanreg
## * For info on the priors used see ?prior_summary.stanreg
```

plot(probit_binom)



So what do the coefficients mean in this model?

```
log_binom$coefficients
```

```
##      PriPos1st Base      PriPos2nd Base      PriPos3rd Base      PriPosCatcher
##      -1.051951        -1.068610        -1.020217        -1.112554
## PriPosCenter Field      PriPosLeft Field      PriPosPitcher      PriPosRight Field
##      -1.028060        -1.050782        -1.906119        -1.027139
##      PriPosShortstop
##      -1.070992
```

```
probit_binom$coefficients
```

```
##      PriPos1st Base      PriPos2nd Base      PriPos3rd Base      PriPosCatcher
##      -0.6469227        -0.6568591        -0.6277234        -0.6828566
## PriPosCenter Field      PriPosLeft Field      PriPosPitcher      PriPosRight Field
##      -0.6326863        -0.6464448        -1.1309219        -0.6325274
##      PriPosShortstop
##      -0.6584987
```

```
invlogit(log_binom$coefficients) * 1000
```

```
##      PriPos1st Base      PriPos2nd Base      PriPos3rd Base      PriPosCatcher
##      258.8506        255.6674        264.9851        247.3950
## PriPosCenter Field      PriPosLeft Field      PriPosPitcher      PriPosRight Field
##      263.4604        259.0749        129.4175        263.6391
##      PriPosShortstop
##      255.2145
```

```
pnorm(probit_binom$coefficients) * 1000
```

```
##      PriPos1st Base      PriPos2nd Base      PriPos3rd Base      PriPosCatcher
##      258.8410        255.6358        265.0926        247.3487
## PriPosCenter Field      PriPosLeft Field      PriPosPitcher      PriPosRight Field
##      263.4692        258.9957        129.0440        263.5211
##      PriPosShortstop
##      255.1089
```